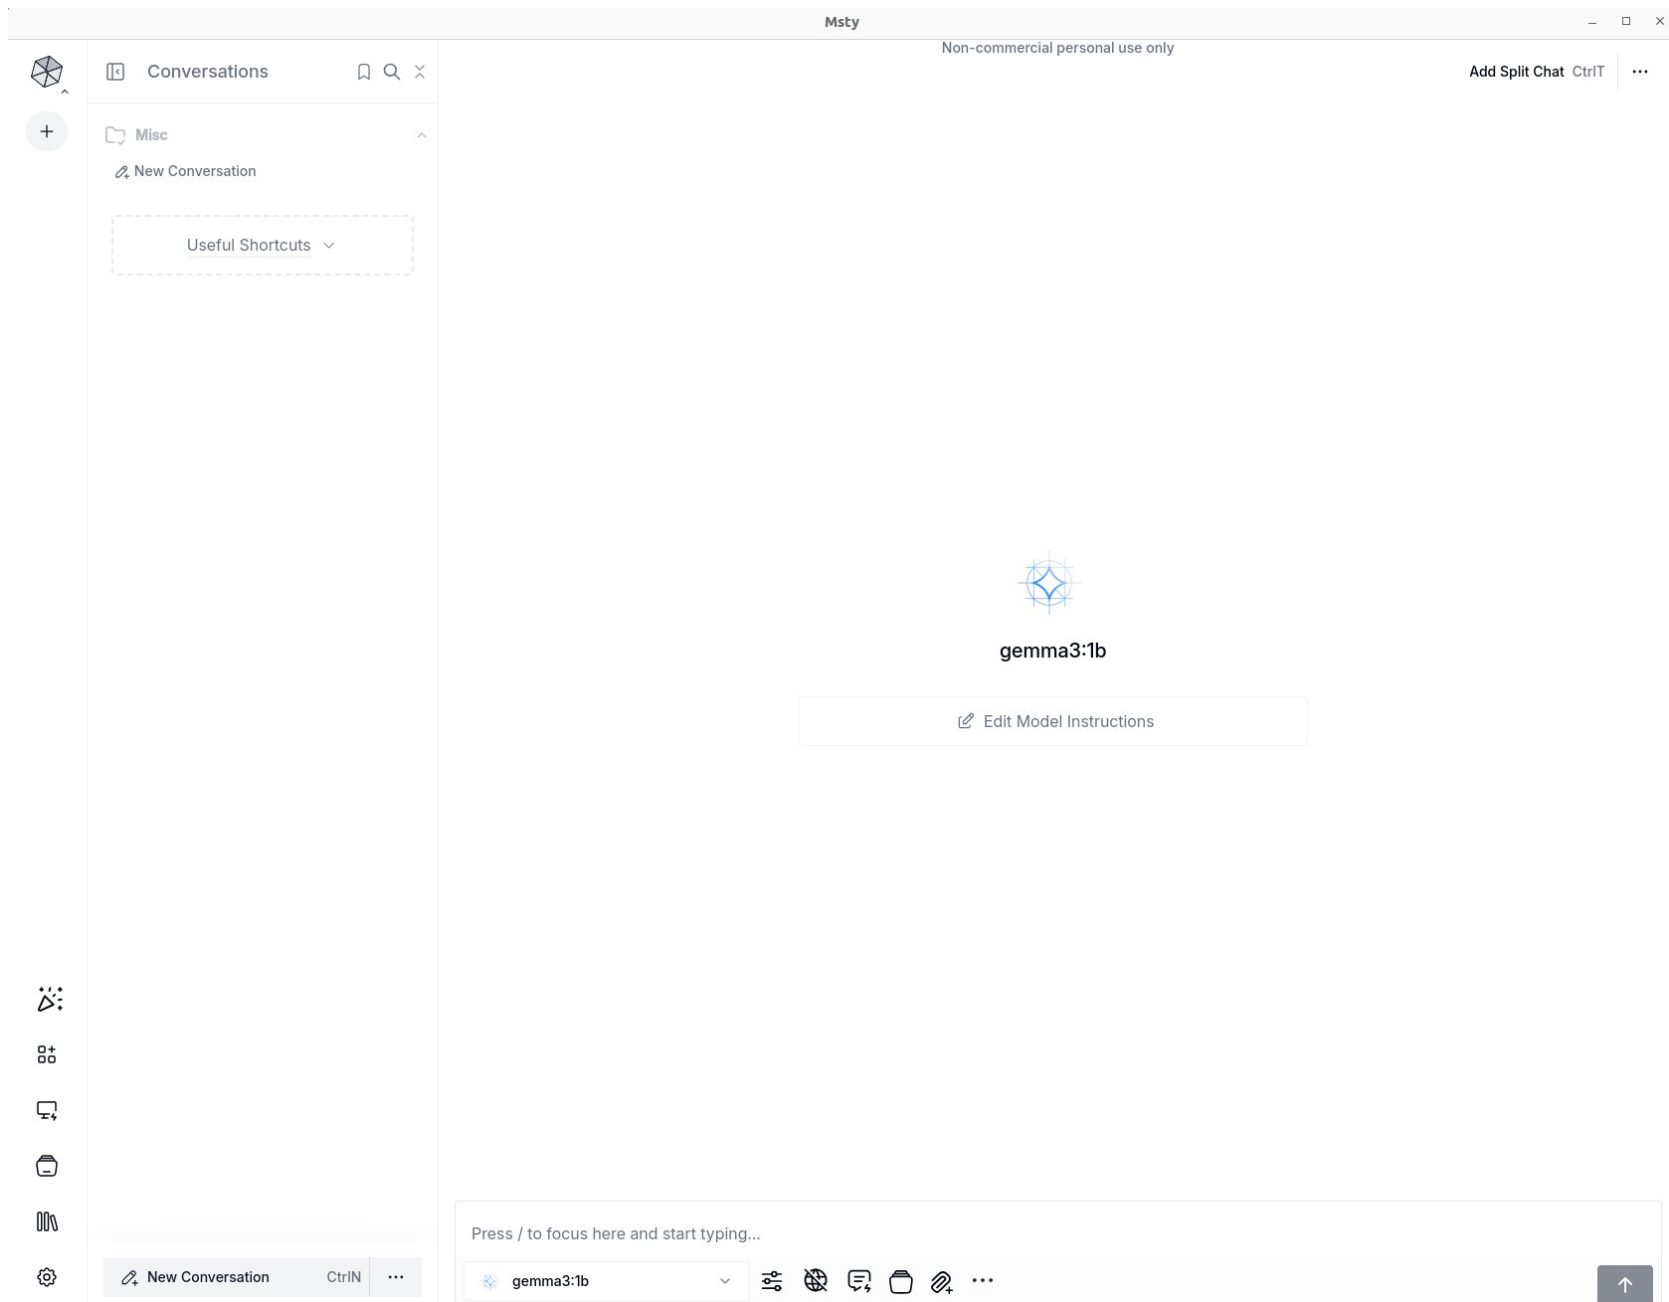
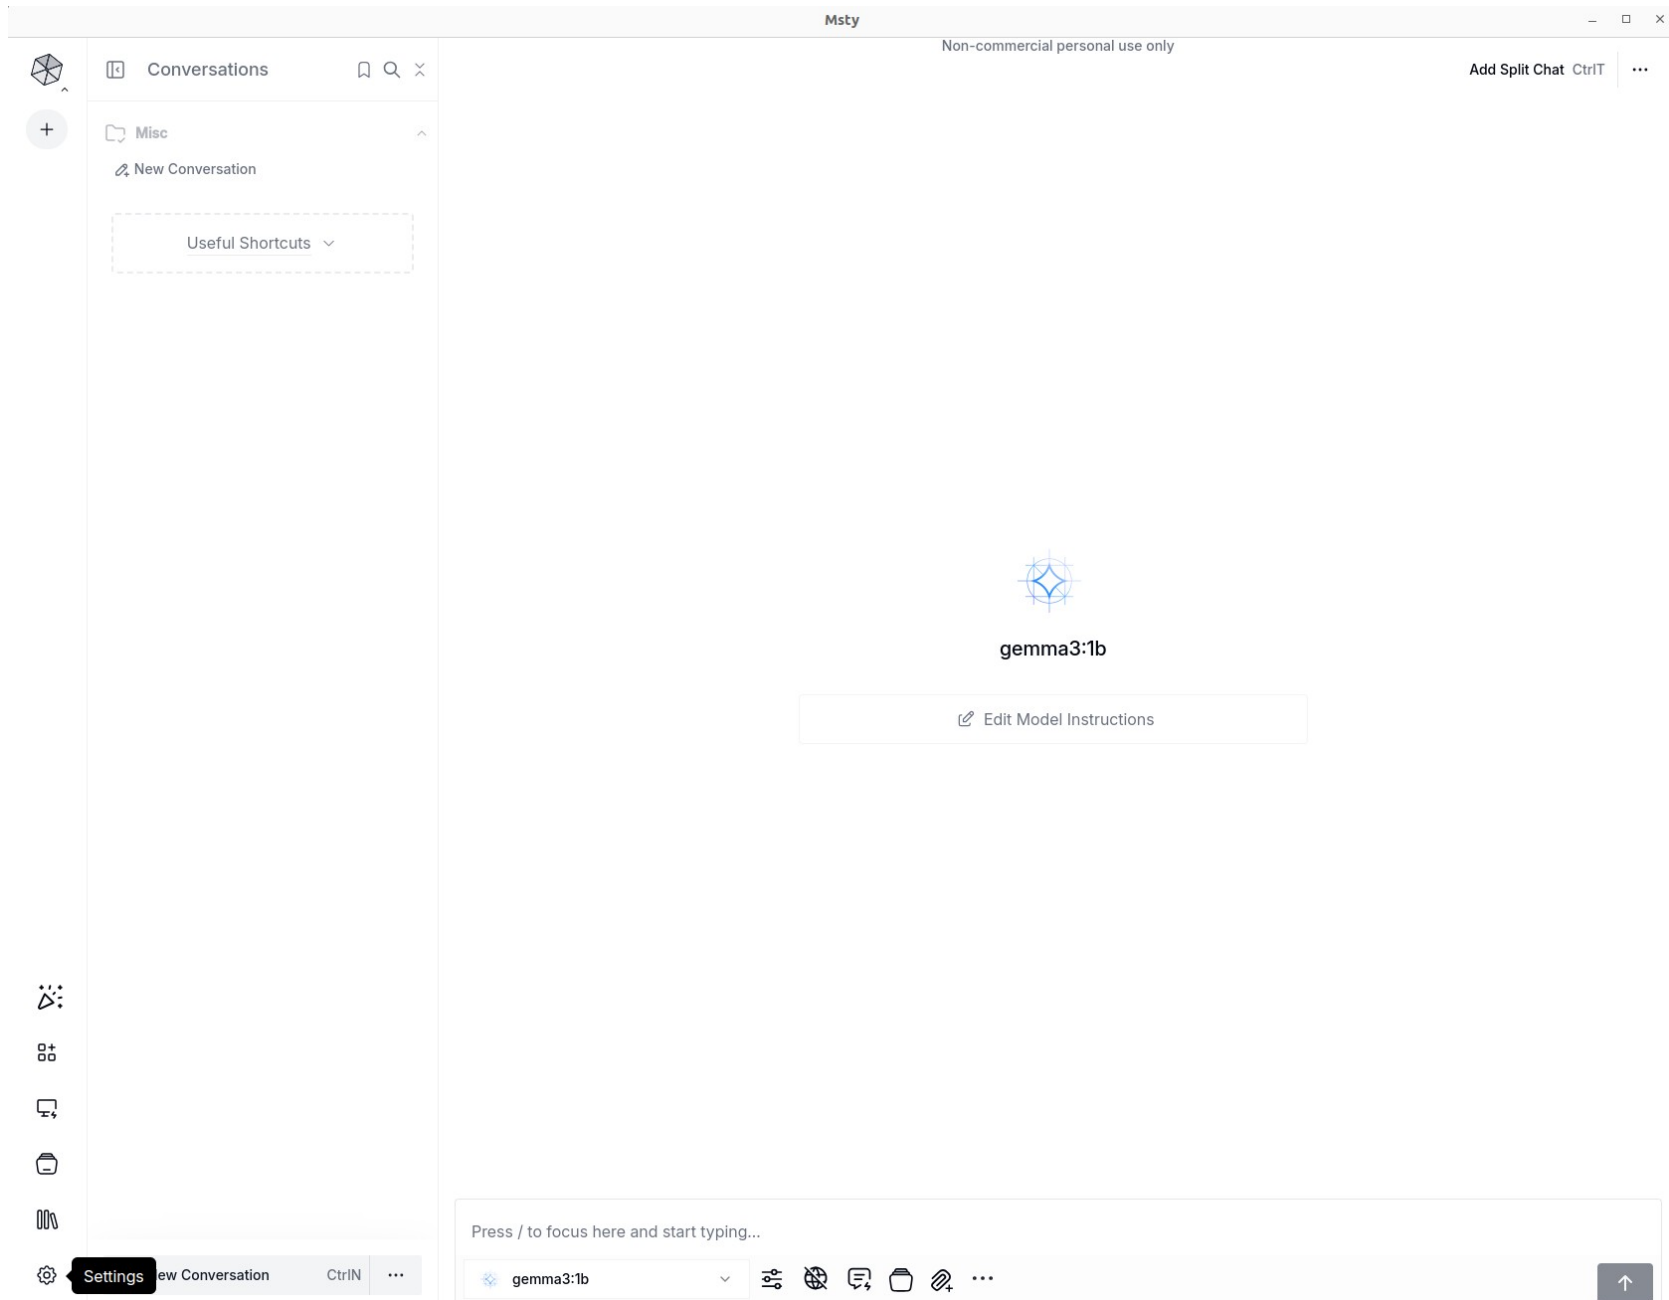


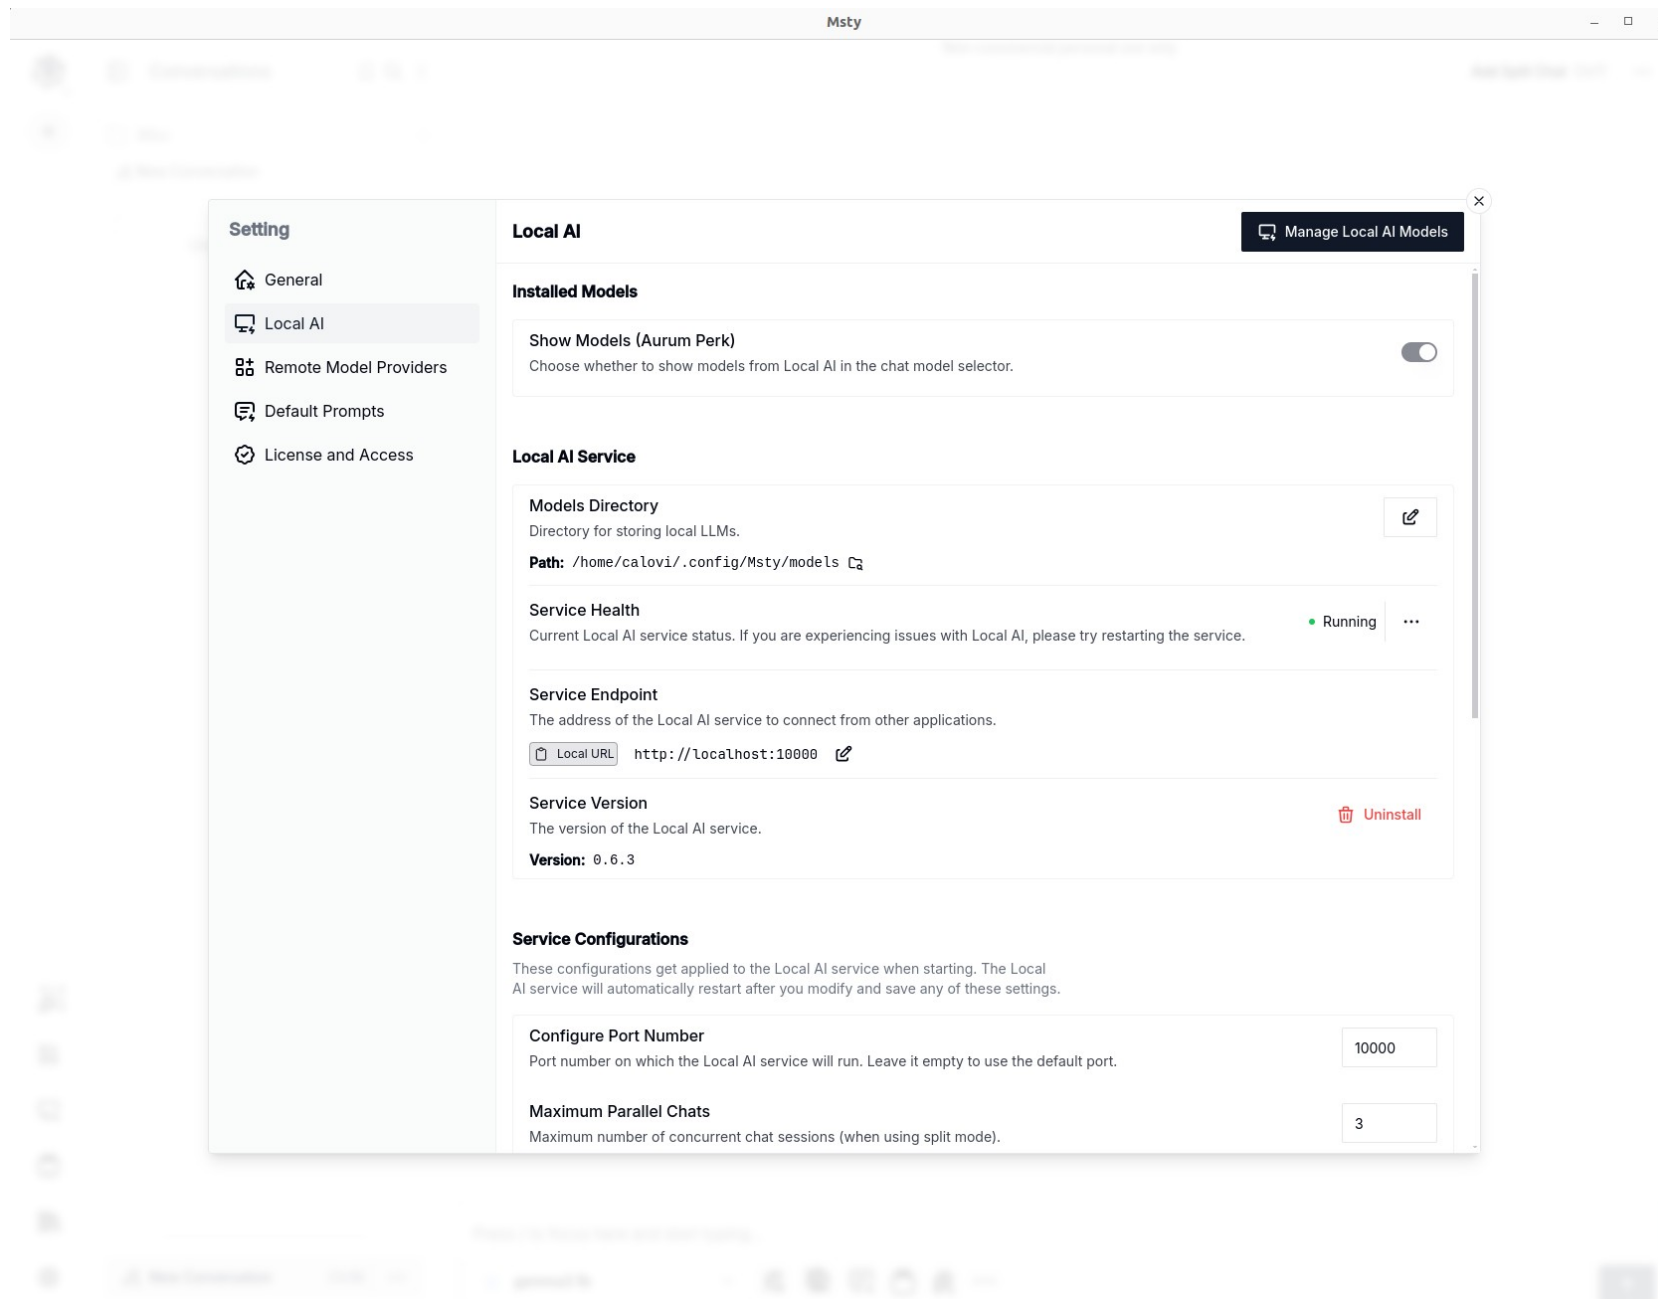
After installing you should be seeing a screen like this



First things first, go to settings



- Then Local AI, and click on Models Directory
 - Models take a lot of space, if you want to save them in a different location, you should change this before you start



Setting

- General
- Local AI
- Remote Model Providers
- Default Prompts
- License and Privacy

Local AI

Installed Models

Show Models (Current Page)
Press enter to show models from local AI or the third-party website

Local AI Service

Model Provider

Local

Web

Server

Cloud

Local

Web

Server

Cloud

Local

Web

Server

Cloud

Service Configuration

These configurations get applied to the local AI service when starting. The local AI service will automatically restart after you modify any of these settings.

Configure Port Number

Port number to which the local AI service will run. Default is 8080.

Maximum Parallel Tasks

Maximum number of concurrent tasks executed using local model.

Edit Models Path

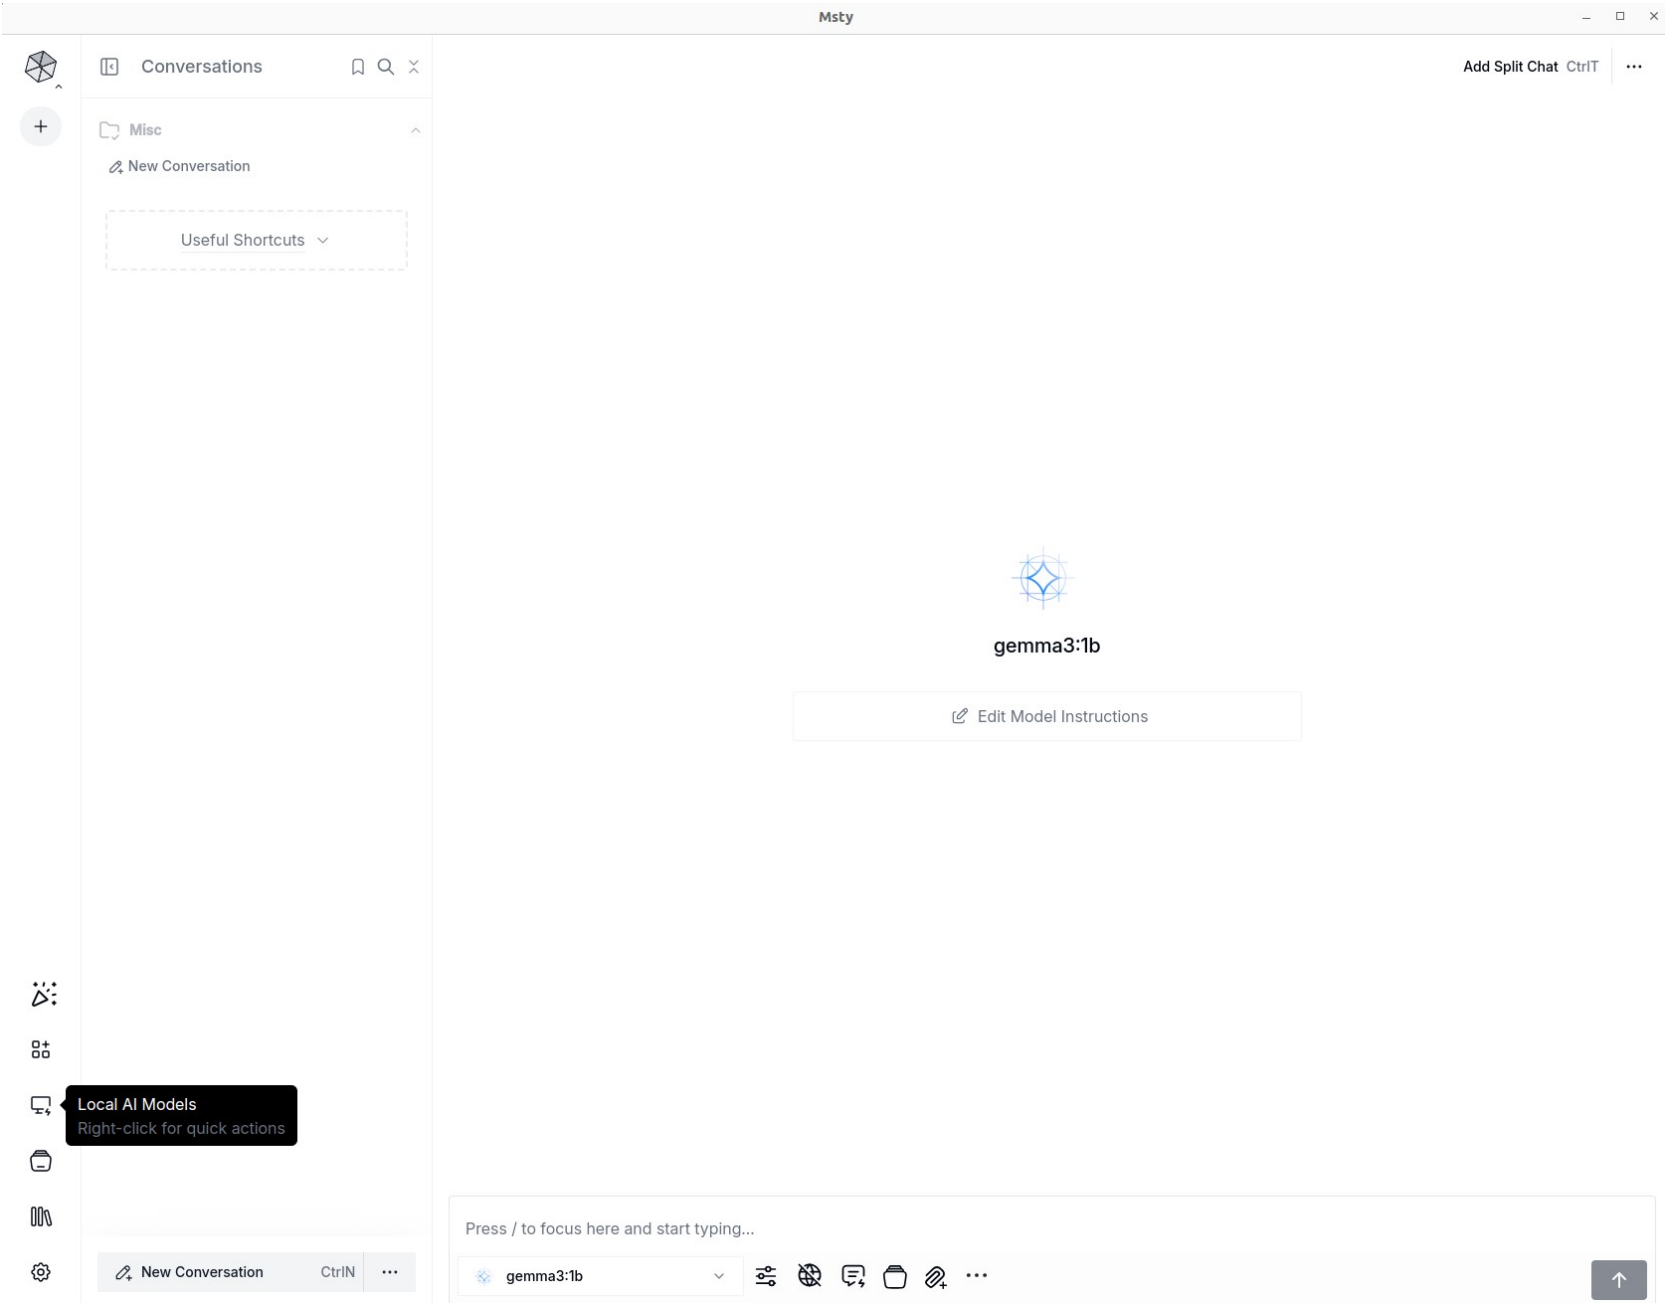
Set your desired location for downloading the models.



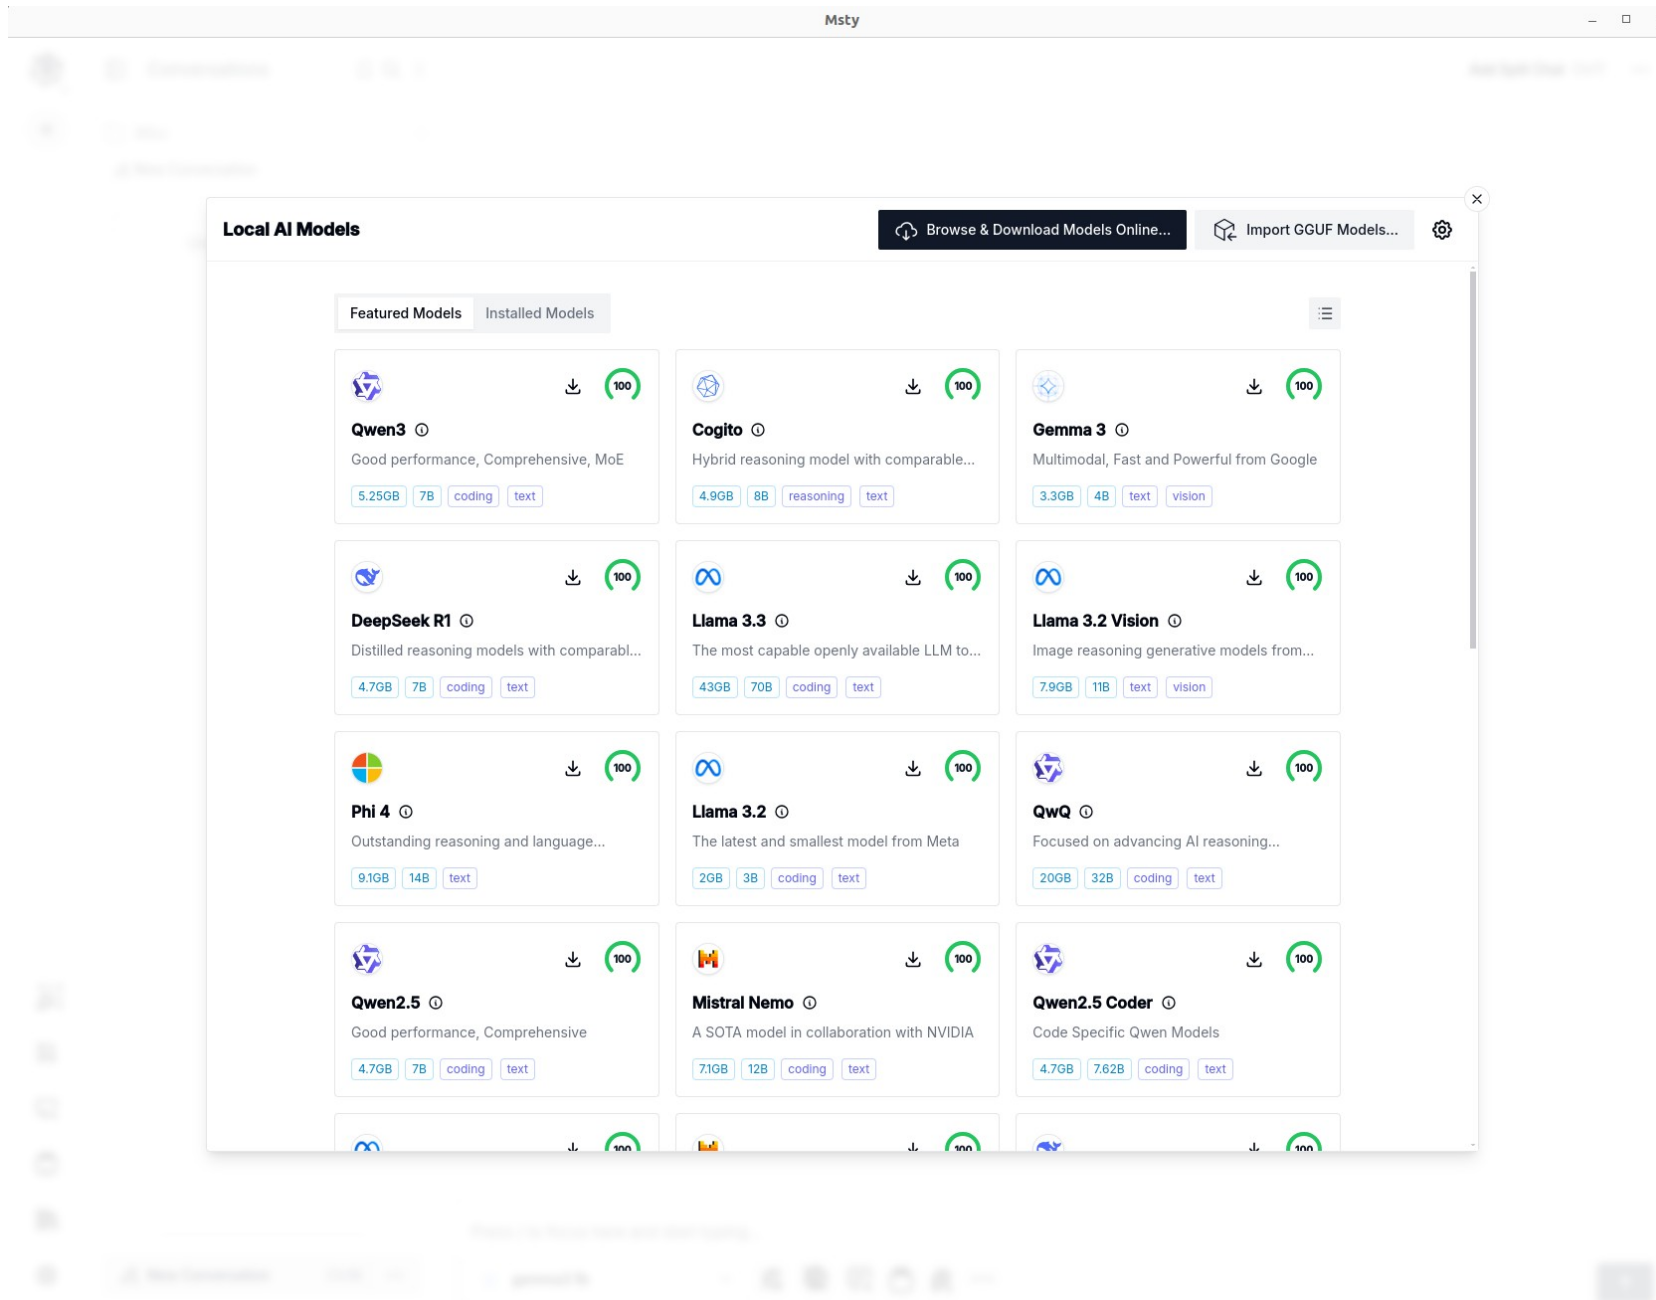
Note: Changing the models path will restart the Local AI service for the changes to take effect.

CancelConfirm

To download models go to Local AI Models



And you can select models to download, notice the ~40GB file size of each model



Local AI Models

[Browse & Download Models Online...](#)[Import GGUF Models...](#)**deepseek-r1:latest**

Installing deepseek-r1:latest... 17.64 %

Featured Models

Installed Models

**Qwen3** ⓘ

Good performance, Comprehensive, MoE

5.25GB 7B coding text

**Cogito** ⓘ

Hybrid reasoning model with comparable...

4.9GB 8B reasoning text

**Gemma 3** ⓘ

Multimodal, Fast and Powerful from Google

3.3GB 4B text vision

**DeepSeek R1** ⓘ

Distilled reasoning models with comparabl...

Installing deepseek-r1:latest... 17.64 %

**Llama 3.3** ⓘ

The most capable openly available LLM to...

43GB 70B coding text

**Llama 3.2 Vision** ⓘ

Image reasoning generative models from...

7.9GB 11B text vision

**Phi 4** ⓘ

Outstanding reasoning and language...

9.1GB 14B text

**Llama 3.2** ⓘ

The latest and smallest model from Meta

2GB 3B coding text

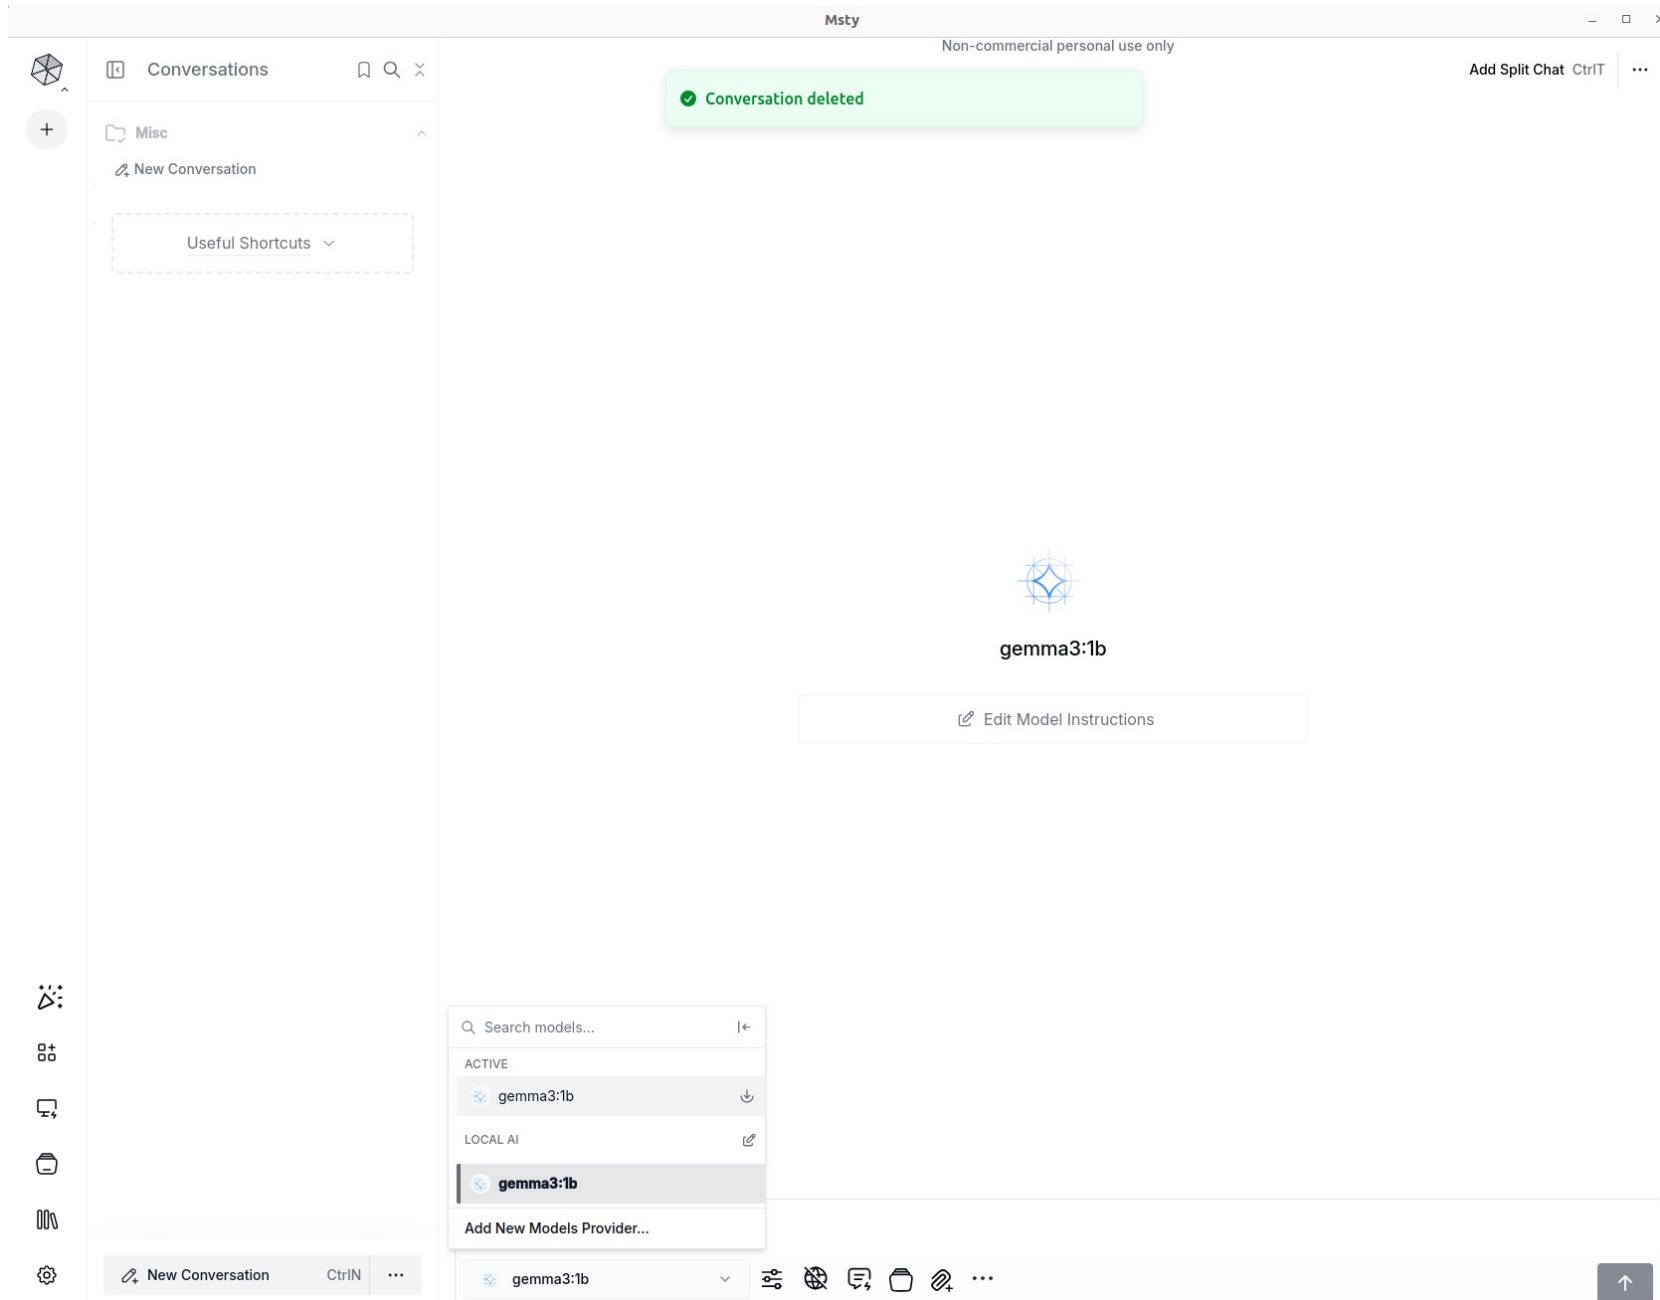
**QwQ** ⓘ

Focused on advancing AI reasoning...


20GB 32B coding text

**Qwen2.5** ⓘ**Mistral Nemo** ⓘ**Qwen2.5 Coder** ⓘ

- Select a Model to use



- Write your prompt









+

Misc


New Conversation

Useful Shortcuts









Conversations


Conversation deleted


gemma3:1b
Edit Model Instructions


Explain to me why quantum entanglement cannot solve FTL communication

gemma3:1b





- Answers will start coming to you



+

Misc

New Conversation

Explain to me why quantum entan...


Useful Shortcuts

Conversations

Model Instructions

You

Explain to me why quantum entanglement cannot solve FTL communication

 **gemma3:1b**

Okay, let's break down why quantum entanglement, while incredibly fascinating, fundamentally can't be used for faster-than-light (FTL) communication. It's a cornerstone of quantum mechanics, and understanding the reasons behind this is crucial. Here's a breakdown of the core issues:

1. The Core of Quantum Entanglement:

- **Linked Fates:** Quantum entanglement is a phenomenon where two or more particles become linked in such a way that they share the same fate, no matter how far apart they are. Measuring a property (like spin or polarization) of one particle instantly tells you the corresponding property of the other, even if they're light-years away.
- **Correlation, Not Communication:** It's important to understand that this isn't about sending a message. The correlation is *inherent* in the quantum state. It's a pre-existing connection.

2. Why It Can't Be Used for FTL Communication – The "Spooky Action"

Here's the critical point: **You can't control the outcome of the measurement.**

- **Randomness Reigns Supreme:** When you measure one entangled particle, you instantly know the state of the other. However, you don't *choose* that state. The outcome is random. Let's say you measure particle A and it's spin-up. The other particle *instantly* becomes spin-down. But, the measurement of particle A has *no effect* on particle B. It's like flipping a coin – you see heads, and you instantly know the other coin is tails, but you didn't *cause* the other coin to be tails.
- **No Signal Transmission:** The problem isn't that the particles are sending information. It's that you can't *manipulate* the measurement outcome to encode a message. Imagine trying to send a message by altering the state of a coin before you flip it. The outcome is random, so you can't decode a signal.

3. Analogy: The Glove Box






A helpful analogy is the glove box:

- **Scenario:** You have a pair of gloves – one left, one right. You randomly separate them and send one to Alice and one to Bob.
- **Classical Communication:** Alice and Bob

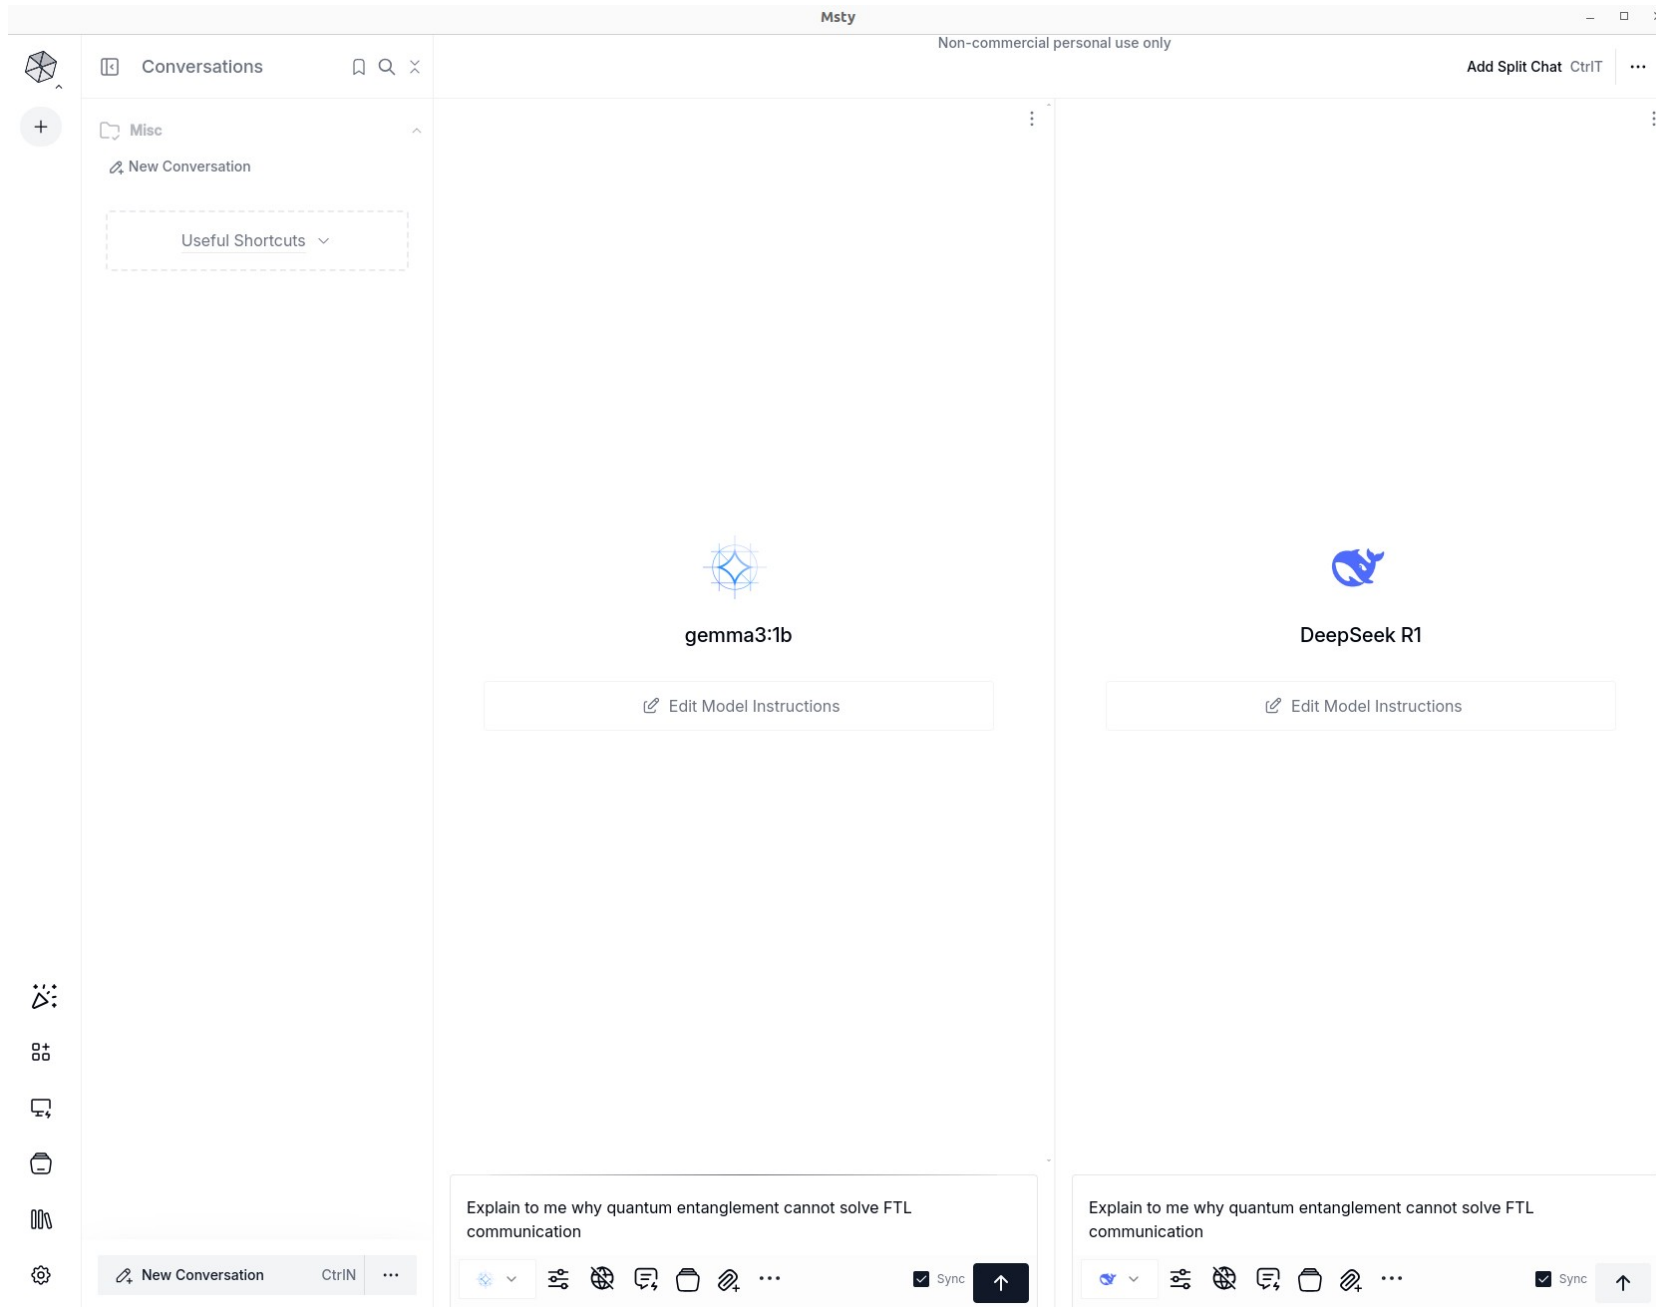
Stop Generation

Press / to focus here and start typing...

gemma3:1b



- You can compare 2 different models with Split chat (top right corner)



- To check the GPU is being used, in Linux you can check your CUDA tools diagnostics by running nvidia-smi

```
(base) calovi@Fisch-XPS:~$ nvidia-smi
Tue Apr 29 17:49:03 2025

+-----+
| NVIDIA-SMI 550.120                Driver Version: 550.120          CUDA Version: 12.4         |
+-----+-----+
| GPU   Name                               Persistence-M | Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp   Perf              Pwr:Usage/Cap |      Memory-Usage | GPU-Util  Compute M. |
|                                           MIG M.         |
+-----+-----+
|  0  NVIDIA GeForce RTX 2060 ...     Off | 00000000:01:00.0 Off |          N/A         |
| N/A   77C    P0               53W /  65W |  4917MiB /  6144MiB |      47%    Default  |
|                                           N/A                 |
+-----+-----+

+-----+
| Processes: |
| GPU   GI    CI          PID    Type   Process name                  GPU Memory |
|      ID    ID              |                 |           Usage         |
+-----+-----+
|  0   N/A   N/A         3635     G   /usr/lib/xorg/Xorg             4MiB      |
|  0   N/A   N/A         9944     C   .../calovi/.config/Msty/bin/msty-local 4908MiB   |
+-----+-----+
```