

# Docker

- Docker is a virtualization software
  - It creates a standalone VM (container) that is able to run cross platform with little efficiency loss (depends on the application, but fair approximation)
- Think of Conda environments, but completely separated from your system
  - Had you ever had to run a software that only runs in Ubuntu xx and uses an outdated package that conflicts with other libraries and softwares?
- Docker containers are versionable
  - Each new addition just gets added to the last version instead of having to redo it

# Learning by doing it

- First thing to check is if docker is running properly in your machine

`docker run hello-world`

```
(base) calovi@Fisch-XPS:~/Dropbox/Konstanz/GPU_CCU/workshop_2024_01/alpine$ docker run hello-world  
Hello from Docker!  
This message shows that your installation appears to be working correctly.
```

After, type

`docker images`

- To delete, type

`docker rmi hello-world`

- Most likely it failed, try:

`docker rmi -f hello-world`

# How to create a docker image?

- Extract folder Alpine into an appropriate location
- Open the file Dockerfile
- Within a terminal, go inside that folder
- Type:  
`docker build -f Dockerfile -t alpine .`

# Running a container

- Type:

```
docker run alpine
```

- Open the file Dockerfile\_2

- What are the differences?
  - What do you imagine it will happen when we run it?

- Now build it and run it:

```
docker build -f Dockerfile_2 -t alpine2 .
```

```
docker run alpine2
```

- Was it what you expected?

# Running a container

- Type:

```
docker run alpine
```

- Open the file Dockerfile\_2

- What are the differences?
  - What do you imagine it will happen when we run it?

- Now build it and run it:

```
docker build -f Dockerfile_2 -t alpine2 .
```

```
docker run alpine2
```

- Was it what you expected?
  - Aren't containers stateless?

# RUN/CMD commands

- **Run** commands are performed during the building of the image
  - The base state of our container already contained the 2 files
- Only the last **CMD** command will be run
  - The first one in the Dockerfile\_2 was ignored

# Stateless version

- Open file Dockerfile\_3, and then build it and run it  
`docker build -f Dockerfile_3 -t alpine3 .`  
`docker run alpine3`
- Note the difference between the **CMD** syntax of the first container to this one
  - Exec form (**CMD** ["command", "param"]) Directly executes commands without a shell, enhancing signal responsiveness and process control
  - Shell form (**CMD** command) Executes commands via a shell, enabling complex scripting such as command chaining and variable expansion
- For complex commands, better to create a bash script and use **CMD** to run it, e.g.  
`CMD ["/usr/local/bin/start-notebook.sh"]`

# Docker is versionable

- Open file Dockerfile\_4, examine it and build it  
`docker build -f Dockerfile_4 -t apine4 .`
- Quite a few more packages were added, now uncomment line 7 (`RUN apk --no-cache add git`) and build the file again
  - Not everything was rebuilt, git was just appended to the image



# Docker is versionable

- Open file Dockerfile\_4, examine it and build it  
`docker build -f Dockerfile_4 -t apine4 .`
- Quite a few more packages were added, now uncomment line 7 (`RUN apk --no-cache add git`) and build the file again
  - Not everything was rebuilt, git was just appended to the image
- Now move line 7 before line 6 and build it again
  - It now had to rebuild package “feh” as well

# Docker is versionable

- When building complex containers it is worth to using **RUN** many times in order to have more saved states
- If building crashes midway through, all completed iterations of **RUN** are already cached
  - imagine having to install opencv multiple times because something at the end of your container crashed?

# Versioning your containers

- In Docker `:latest` is just the default version, not exactly the latest version
- Instead of creating multiple images (alpine, alpine2, alpine3), we can create different versions of them

- Type:

```
docker build -f Dockerfile -t alpine:1.0 -t alpine:latest .
```

```
docker build -f Dockerfile_3 -t alpine:2.0 .
```

```
docker run alpine
```

```
docker run alpine:2.0
```

- Tags can be anything and are case sensitive, `latest` and `Latest` would refer to different versions

# Repositories

- When building an image you can use default packages like
  - FROM alpine:latest
  - FROM quay.io/jupyter/base-notebook
  - FROM nvcr.io/nvidia/tensorflow:21.02-tf2-py3
- But you can also download a pre-packaged container

**This will download a 1GB container**

`docker run -p 8888:8888 jupyter/base-notebook`

- Where the port syntax is: `-p localhost:container`

# More complex example

- Open file Dockerfile\_Jupyter
  - (I did some last minute trimming, might not work from scratch)

# Logging in the container

- You have a container, and you want to get inside that VM, type:

```
docker run -d --name running_alpine alpine tail -f /dev/null
```

- -d is to keep the terminal free (detached)
- --name is to ensure you give an specific name

- Type

```
docker ps
```

```
docker exec -it running_alpine /bin/sh
```

- Now you are actually inside the pod, and while it is active you can perform changes to it

# Uploading your container

- To upload your container you need to be logged in a repository and have permission to **push** (upload) it
- Before pushing, one needs to update the image name to have the address of the repository, i.e.

```
docker tag ccu-workshop-jupyter
```

```
ccu-k8s.inf.uni-konstanz.de:32250/daniel.calovi/ccu-workshop-jupyter
```

```
docker push ccu-k8s.inf.uni-konstanz.de:32250/daniel.calovi/ccu-workshop-jupyter
```

- This will not run for you, just an example for later

# Questions/Lunch Break?