

# What is Kubernetes and why we use it?

# What is Kubernetes and why we use it?

- Automated Container Management
  - Automates deployment, scaling, and operations of application containers across clusters of hosts.
- Ensures Reliability
  - Maintains high availability, facilitates scaling, and manages rollouts and rollbacks.
- Resource Optimization
  - Efficiently manages resource allocation, maximizing utility and minimizing waste

# What is Kubernetes and why we use it?

- Automated Container Management
  - Automates deployment, scaling, and operations of application containers across clusters of hosts.
- Ensures Reliability
  - Maintains high availability, facilitates scaling, and manages rollouts and rollbacks.
- Resource Optimization
  - Efficiently manages resource allocation, maximizing utility and minimizing waste
- All major cloud computing solutions use it

# Who here already tried to use our cluster?

# Step by step guide to our cluster

- After receiving the login details, you will be able to log in here:

<https://ccu-k8s.inf.uni-konstanz.de>

## Welcome to the CCU compute cluster

Note: all services below can only be accessed from the university network (university VPN or a local computer).

### Login services

Use this link to obtain an authentication token for kubectl (see below documentation for details).

[Login to the cluster](#)

If your account has just been created, or if you have forgotten your password, you can use the following link to obtain a password reset email.

[Password reset](#)

### Help and documentation

As a starting point, here is a tutorial on how to get your workloads to the GPU cluster. A lot of further documentation is also linked there.

[Quick start tutorial](#)

**Note: the password database is copied to the Wiki every few minutes. If you changed your password, allow some time to pass before trying to log in to the Wiki.**

General help on how to use Kubernetes:

[kubectl cheat sheet \(do yourself a favor and enable kubectl bash autocompletion\)](#)

### Diagnostics

[Cluster status page](#)

# Wiki

- At the CCU wiki there is a Quick Start guide
  - It will improve significantly in the near future
- The Log in section is well described, steps are basically:
  - Install kubectl in your computer
  - Login in the first link
  - Use the Full Kubeconfig section
  - Alter as described in the wiki
- If everything went properly, type:  
`kubectl config use-context ccu-k8s`  
`kubectl get pods`
  - Answer will be: No resources found in namespace user-your-name

# Repository access (soon in the wiki)

- Once your kubectl is configured, you should login to the cluster repository:

```
sudo docker login ccu-k8s.inf.uni-konstanz.de:32250
```

- Check that you have a file at:

```
/home/calovi/.docker/config.json
```

- Create a **secret**

```
kubectl create secret generic YOURTAG --from-  
file=.dockerconfigjson=/home/calovi/.docker/config.json --  
type=kubernetes.io/dockerconfigjson
```

- The **YOURTAG** is what is going to be used in every script submission to the cluster

# Pods and Jobs

- When submitting a script to the cluster one usually submit a Pod, or a Job
  - Pod is a continually running container, like the jupyter notebook server, or a session that you want to get into the VM and test what it works
  - Job, is just a sequence of commands, that is run, and once it is finished it automatically closes  
(having the care to save your data in your permanent storage)
- Pods once created, will remain active, and using resources until deleted



# Best practices

- Whenever testing a new container use Pods. Test it, run it, check that everything is running well
- Delete your pod to free resources for the rest of the community
- You will receive a permanent storage in the cluster under your namespace (login), the rule of thumb is:
  - Large files should be stored in the permanent storage
  - Many small training files, should be copied to the pod/job temporary storage

# Best Practices

- Do you have a data management plan already?
  - If so, soon we will have tutorials how to mount your hot backup drive directly into the cluster



Ilja Werner

Don't forget to create a data management plan



# Best Practices

- Do you have a data management plan already?
  - If so, soon we will have tutorials how to mount your hot backup drive directly into the cluster
- Your storage **WILL** fail eventually, better be prepared



Ilja Werner

Don't forget to create a data management plan



# Ok, but what is Kubernetes and how do we use it?

- A typical script for a pod will look like this
- Don't get intimidated, most of the things you don't have to change

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-test-pod
spec:
  containers:
  - name: ubuntu
    image: ubuntu:20.04
    command: ["sleep", "1d"]
    resources:
      requests:
        cpu: 100m
        memory: 100Mi
      limits:
        cpu: 1
        memory: 1Gi
    volumeMounts:
      - mountPath: /abyss/home
        name: cephfs-home
        readOnly: false
      - mountPath: /abyss/shared
        name: cephfs-shared
        readOnly: false
      - mountPath: /abyss/datasets
        name: cephfs-datasets
        readOnly: true
  volumes:
  - name: cephfs-home
    hostPath:
      path: "/cephfs/abyss/home/<your-username>"
      type: Directory
  - name: cephfs-shared
    hostPath:
      path: "/cephfs/abyss/shared"
      type: Directory
  - name: cephfs-datasets
    hostPath:
      path: "/cephfs/abyss/datasets"
      type: Directory
```

# Kubectl cheat sheet

- `kubectl apply -f Your_Script.yaml`
  - How to submit your pod or job to the cluster
- `kubectl get pods`
  - Describe all your running pods
- `kubectl describe pods Your_Script`
  - Shows the basic status of your pod and possibly why it failed to be mounted
- `kubectl log pods/Your_Pod`
  - Terminal messages of your pod, possible errors and other messages
- `kubectl delete pods Your_Pod`
- `kubectl delete -f Your_script.yaml`
  - If using the script option, kills all nested services
  - Deleting your Pod from the server
    - Deletes everything that was not saved in your permanent storage
- `Kubectl exec -it Your_Pod /bin/bash`
  - Same as docker, get inside a running pod

# Cluster nodes

- The cluster has different nodes
- Each node can have a different type of GPU, with different attributes

CCU name	Access	Platform	GPUs	Labels	Taints
<b>imp</b>	all	Dual Xeon Rack	4 x Titan Xp @ 12 GB	gpumem=12, gpuarch=nvidia-titan, nvidia-compute-capability-sm70=true	
<b>dretch</b>	all	Dual Xeon Rack	4 x Titan RTX @ 24 GB	gpumem=24, gpuarch=nvidia-titan, nvidia-compute-capability-sm70=true	
<b>belial</b>	exc-cb	Supermicro	8 x Quadro RTX 6000 @ 24 GB	gpumem=24, gpuarch=nvidia-rtx, nvidia-compute-capability-sm75=true	gpumem=24:NoSchedule
<b>fierna</b>	exc-cb	Supermicro	8 x Quadro RTX 6000 @ 24 GB	gpumem=24, gpuarch=nvidia-rtx, nvidia-compute-capability-sm75=true	gpumem=24:NoSchedule
<b>vecna</b>	exc-cb, inf	nVidia DGX-2	16 x V100 @ 32 GB	gpumem=32, gpuarch=nvidia-v100, nvidia-compute-capability-sm80=true	gpumem=32:NoSchedule
<b>zariel</b>	trr161	nVidia DGX A100	8 x A100 @ 40 GB	gpumem=40, gpuarch=nvidia-a100, nvidia-compute-capability-sm80=true	gpumem=40:NoSchedule
<b>tiamat</b>	exc-cb	Supermicro	4 x A100 @ 40 GB	gpumem=40, gpuarch=nvidia-a100, nvidia-compute-capability-sm80=true	gpumem=40:NoSchedule
<b>asmodeus</b>	all	Supermicro	4 x A100 HGX 320 GB, subdivided in 8 GPUs @ 40 GB	gpumem=40, gpuarch=nvidia-a100, nvidia-compute-capability-sm80=true	gpumem=40:NoSchedule
<b>demogorgon</b>	exc-cb	Delta	8 x A40 @ 48 GB	gpumem=48, gpuarch=nvidia-a40, nvidia-compute-capability-sm80=true	gpumem=48:NoSchedule

# Taints

- To access a specific node, one can do it by describing correct “**taint**”
  - Detailed info at the wiki
- One can select a node by name as well by using the **nodeSelector**

# Port forwarding

- At the moment, to have access to the graphical interface inside a pod, one would have to access it via `ssh -X`
  - Better solutions are being developed, but were not ready for today
- In the `Dockerfile_Jupyter` one can see how I managed to install `ssh` and to have access to it
- The container enables the usual **port 22** for `ssh` within the container, but to access it, one has to create a rule for forwarding the port to our local machine

```
kubectl port-forward <pod-name> <dest-port>:<source-port>
```
- If everything was set up appropriately, one can then access it by:
  - `ssh -X -p <dest-port> root@localhost`
- More commonly this would be used to access browser related pods like jupyter-notebook servers



# Port forwarding

- **WARNING**
  - For most situations, graphical interface to the pod is complete unnecessary
  - Labelling data should be done at your local machine and files transferred to your permanent/temporary storage
  - You can be putting your machine and the cluster at risk depending on how secure your machine is  
(unlikely, but good practice to be cautious)
- `ssh -X` is a VERY slow way to handle visual interface, if needed check with Ilja and me to see if our alternative is already in place

# Questions, break?